


Towards Interactive High-Performance Computing with ALCF JupyterHub

Murat Keçeli

keceli@anl.gov

Computational Science Division, Argonne National Laboratory



Outline

- Project Jupyter
- What you can do with Jupyter?
- Jupyter/IPython basics
- Introduction to markdown, magic, widgets
- Introduction to ALCF JupyterHub
- Live Demos
 - New kernel installation
 - ezCobalt: how to submit jobs
 - ezBalsam: how to use Balsam

Disclaimer

- This webinar will not cover:
 - low level details about queuing or ensembling jobs or creating Balsam workflows, etc. covered in a [previous webinar](#)
 - using Jupyter through an ssh tunnel, reverse proxy, or remote kernels
 - using Dask, Spark, Kubernetes, or a container for distributed computing
 - accessing compute nodes directly
- ALCF JupyterHub is a new service and improving rapidly. You can send an email to support@alcf.anl.gov (cc: keceli@anl.gov) for problems and suggestions.

Project Jupyter

- Started in 2014, as an IPython spin-off project led by Fernando Perez to “develop open-source software, open-standards, and services for interactive computing”.
- Inspired by Galileo’s notebooks and languages used in scientific software: Julia, Python, and R.

Observations: $\frac{1}{2}$ in

2. $\frac{1}{2}$ in
March 19. $\bigcirc \times \times$

5. $\frac{1}{2}$ in
April. $\times \times \bigcirc \times$

7. $\frac{1}{2}$ in
April. $\bigcirc \times \times \times$

3. $\frac{1}{2}$ in
May. $\bigcirc \times \times$

3. $\frac{1}{2}$ in
May. $\times \bigcirc \times$

4. $\frac{1}{2}$ in
May. $\times \bigcirc \times \times$

6. $\frac{1}{2}$ in
May. $\times \times \bigcirc \times$

8. $\frac{1}{2}$ in
May 17. $\times \times \times \bigcirc$

10. $\frac{1}{2}$ in
May. $\times \times \times \bigcirc \times$

11. $\frac{1}{2}$ in
May. $\times \times \bigcirc \times$

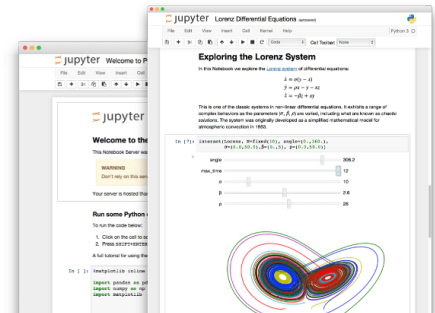
12. $\frac{1}{2}$ in
May. $\times \times \bigcirc \times$

14. $\frac{1}{2}$ in
May. $\times \times \times \times$

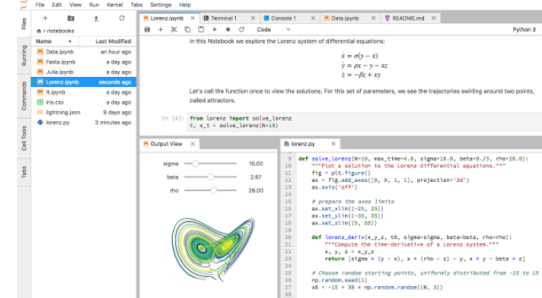


Jupyter X

Jupyter Notebook



JupyterLab



jupyter {book}

What you can do?

- Interactive development environment
 - Fast code prototyping, test new ideas easily
 - Most languages are supported through [Jupyter kernels](#)
- Learn or teach with notebooks
 - Prepare tutorials, run demos
- Data analysis and visualization
- Presentations with Reveal.js
- Interactive work on HPC centers or cloud
 - JupyterHub
 - [Google Colab](#)
 - [Binder](#)

Basics (Shortcuts)

- `Esc/Enter` get in command/edit mode

Command mode	Edit mode
<code>h</code> show (edit) all shortcuts	<code>shift enter</code> Run cell, select below
<code>a/b</code> insert cell above/below	<code>cmd/ctrl enter</code> Run cell
<code>c/x</code> copy/cut selected cell	<code>tab</code> completion or indent
<code>V/v</code> paste cell above/below	<code>shift tab</code> tooltip
<code>d,d</code> delete cell	<code>cmd/ctrl d</code> delete line
<code>y/m/r</code> code/markdown/raw mode	<code>cmd/ctrl a</code> select all
<code>f</code> search, replace	<code>cmd/ctrl z</code> undo
<code>p</code> open the command palette	<code>cmd/ctrl /</code> comment

Basics (Shortcuts)

- `Esc/Enter` get in command/edit mode

Command mode	Edit mode
<code>h</code> show (edit) all shortcuts	<code>shift enter</code> Run cell, select below
<code>a/b</code> insert cell above/below	<code>cmd/ctrl enter</code> Run cell
<code>c/x</code> copy/cut selected cell	<code>tab</code> completion or indent
<code>V/v</code> paste cell above/below	<code>shift tab</code> tooltip
<code>d,d</code> delete cell	<code>cmd/ctrl d</code> delete line
<code>y/m/r</code> code/markdown/raw mode	<code>cmd/ctrl a</code> select all
<code>f</code> search, replace	<code>cmd/ctrl z</code> undo
<code>p</code> open the command palette	<code>cmd/ctrl /</code> comment

In [6]:

```
import os
os.getenv??
#help('modules')
#help('modules mpi4py')
```


Markdown

- bullet list
 - subbullet
- equation: $E = mc^2$
- inline code `echo hello jupyter``
- A [link](#)
- Table

Col 1	Col 2	Col 3
1, 1	1,2	1,3
2, 1	2,2	2,3
3, 1	3,2	3,3

- A kitten



IPython Magic

- Magic functions are prefixed by `%` (line magic) or `%%` (cell magic)
- Cell magic `%%` should be at the first line
- Shell commands are prefixed by `!`
- `%quickref` : Quick reference card for IPython
- `%magic` : Info on IPython magic functions
- `%debug` : Interactive debugger
- `%timeit` : Report time execution
- `%prun` : Profile (`%lprun` is better, `pip install lprun` and `%load_ext line_profiler`)

```
In [7]: %magic
```

```
In [7]: %magic
```

```
In [60]: import numpy as np
a = [1]*1000
%timeit sum(a)
b = np.array(a)
%timeit np.sum(a)
%timeit np.sum(b)
```

10000 loops, best of 5: 7.51 μ s per loop

The slowest run took 145.08 times longer than the fastest. This could mean that an intermediate result is being cached.

10000 loops, best of 5: 106 μ s per loop

The slowest run took 5.23 times longer than the fastest. This could mean that an intermediate result is being cached.

100000 loops, best of 5: 7.14 μ s per loop

Jupyter Widgets (ipywidgets)

- Widgets are basic GUI elements that can enhance interactivity on a Jupyter notebook
- Enables using sliders, text boxes, buttons, and more that can link input and output.

Jupyter Widgets (ipywidgets)

- Widgets are basic GUI elements that can enhance interactivity on a Jupyter notebook
- Enables using sliders, text boxes, buttons, and more that can link input and output.

In [1]:

```
import ipywidgets  
ipywidgets.IntSlider()
```

Jupyter Widgets (ipywidgets)

- Widgets are basic GUI elements that can enhance interactivity on a Jupyter notebook
- Enables using sliders, text boxes, buttons, and more that can link input and output.

```
In [1]: import ipywidgets  
        ipywidgets.IntSlider()
```

```
In [2]: ipywidgets.Text(value='Hello Jupyter!', disabled=False)
```

Jupyter Widgets (ipywidgets)

- Widgets are basic GUI elements that can enhance interactivity on a Jupyter notebook
- Enables using sliders, text boxes, buttons, and more that can link input and output.

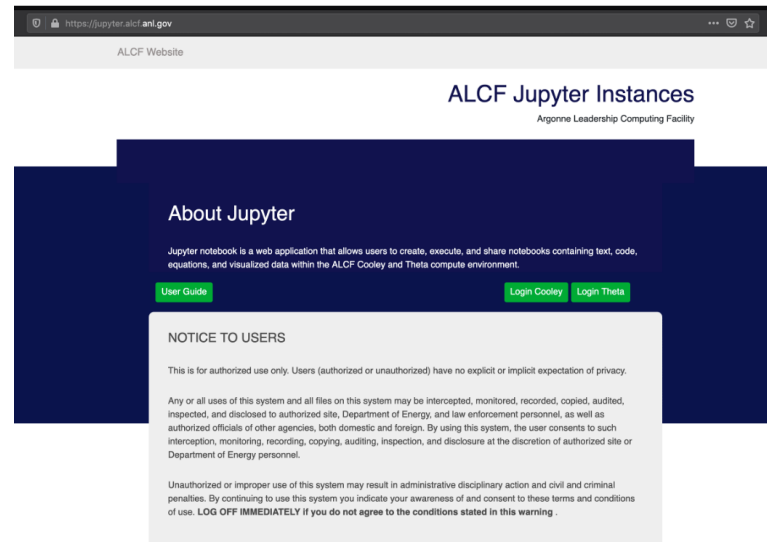
```
In [1]: import ipywidgets  
        ipywidgets.IntSlider()
```

```
In [2]: ipywidgets.Text(value='Hello Jupyter!', disabled=False)
```

```
In [3]: ipywidgets.ToggleButton(value=False, description="Don't click",  
                                button_style='danger', tooltip='Description',)
```


ALCF JupyterHub

- If you are an ALCF user, you can log in to Jupyter Hub at <https://jupyter.alcf.anl.gov> using your ALCF credentials.
- If not, check <https://alcf.anl.gov/support-center/get-started>
- Jupyter Hub instances runs on an external servers, but not on login, mom, or compute nodes.
- Servers have 16 core Intel(R) Xeon(R) CPU E5-2683 and 512 GB memory and reserved for data analytics and visualization, not simulations.



Sign in

Username:

Password:

Sign In

ALCF JupyterHub

- JupyterHub for Cooley :
 - runs on `jupyter01.mcp.alcf.anl.gov`
 - has access to the user's home folder `/home/$USER`, the Mira projects folder `/projects`, and the Theta project folder `/lus/theta-fs0/projects`
 - submitted jobs will run on Cooley
- JupyterHub for Theta:
 - runs on `jupyter02.mcp.alcf.anl.gov`
 - has access to your home folder `/home/$USER` and projects folder `/lus/theta-fs0/projects`*
 - does not have access to `/opt/cray`, `/opt/intel`, etc., that is, you cannot use any Theta modules or any Cray libraries.
 - Submitted jobs will run on Theta

Notes

- JupyterHub starts on your home folder, to access project folders, you can create a symbolic link `!ln -s /project/my_project my_project`
- If you have a broken symlink on your home directory, JupyterHub gives a server error with `permission denied` message. You need to clean up / fix the broken symbolink links.
- When you exceed your file quota, you may also experience problems. Check with `myquota`.
- To run JupyterLab on JupyterHub, modify the link to `https://jupyter.alcf.anl.gov/cooley/user/$USER/lab`
- Documentation is available at <https://www.alcf.anl.gov/support-center/theta/jupyter-hub>

How to install a new Conda environment & Jupyter kernel

Step 0

- Check the names of the existing environments & kernels:

```
!conda env list  
!jupyter kernelspec list
```

- Select a name for the new environment & kernel.
- Using a prefix such as `jhub_` is helpful to distinguish JupyterHub environments from others.

```
ENVNAME="jhub_demo"
```

Step 1

- Create a new environment

```
!conda create -y --name $ENVNAME
```

- Or, create a new environment with a different python version

```
!conda create -y --name $ENVNAME python=3.8
```

- Or, create a new environment with a clone of the base environment (recommended)

```
!conda create -y --name $ENVNAME --clone base
```

A step backward

Step 2

- Install new packages with `conda`, or `pip`

```
!source activate $ENVNAME; conda install -y -c conda-forge rise
```

```
!source activate $ENVNAME; pip install balsam-flow
```

- Or, if you didn't clone from the base, you need to install the following packages additionally:

```
!source activate $ENVNAME; conda install -y jupyter nb_conda  
ipykernel
```

Step 3

- Install the kernel for Jupyter

```
!source activate $ENVNAME;python -m ipykernel install --user --name  
$ENVNAME
```


Final steps

- Refresh the browser or open a new notebook.
- Select the new `Kernel` from the top dropdownlist
- When you need to install another package, you only need to run the following steps

```
ENVNAME='jhub_demo'  
!source activate $ENVNAME; conda install -y <any_conda_package>  
!source activate $ENVNAME; pip install -c <any_pypi_package>
```

Notes

- Check the installation with

```
!conda list
import <any_package>
print(<any_package>.__file__)
print(<any_package>.__version__)
```

- Do not use environments installed on JupyterHub elsewhere.

Clean up

- You may run out of space quickly, check with `myquota`.
- You can run `conda clean` to remove index cache, lock files, tarballs, unused cache packages, and source cache

```
!conda clean --all -y
```

- To remove an environment and the kernel you don't need:

```
!conda env remove -y -n $ENVNAME  
!jupyter kernelspec uninstall -y $ENVNAME
```

Resources

- [Fernando Perez's Project Jupyter presentation](#)
- [jupyter.org](#)
 - Check out Voilà, Jupyter Lab, Jupyter Book
- [Jupyter tutorial](#)
- [Version control for Jupyter](#)
- [ALCF ML tutorials](#)
- [More ALCF notebooks](#)



Acknowledgements



- Thank you all for attending
- Thanks to Misha, Alvaro, and Ray for their feedback and suggestions
- Thanks to Tommie for running and maintaining JupyterHub servers
- Thanks to Gurunath for working together during the summer
- Thanks to Venkat, Tom, and Mike for motivation and support

Live Demo

- All materials are at <https://github.com/keceli/ezHPC>

```
git clone https://github.com/keceli/ezHPC
```

